

Self-Test

SER 50 R Workshop

S Mooney

June 20, 2017

1. Use `c` to create a vector named `myNumbers` from these numbers:
9 9 5 8 4 2 8

```
> myNumbers <- c(9, 9, 5, 8, 4, 2, 8)
>
```

2. Create a vector named `myLetters` of 26 lowercase letters from a to z.
(Hint: check out the `letters` vector is built into R)

```
> myLetters <- letters
```

3. Combine the vectors `myNumbers` and `myLetters` columnwise into a matrix called `myObj`.

```
> myObj <- cbind(myNumbers, myLetters)
```

4. Create a vector of numbers named `myvec` with the numbers from 19 to 10 (i.e. 19, then 18, then 17.)

```
> myvec <- 19:10
```

5. Sort `myvec` in ascending order

```
> myvec <- sort(myvec)
```

6. Create a logical vector called `teens` with the value `TRUE` for all numbers in `myvec` that are more than 12 (i.e. your resulting vector should start `FALSE FALSE FALSE TRUE TRUE`.)

```
> teens <- myvec > 12
```

7. Create an object containing one item named `'mynumbers'` with values 1,2,3, and 4, and one item named `'myletters'` and has values e, f, and g.
(Hint: what basic object type should this be?)

```
> list(mynumbers=1:4, myletters=c('e', 'f', 'g'))
```

```
$mynumbers
[1] 1 2 3 4
```

```
$myletters
[1] "e" "f" "g"
```

8. Run the following code:

```
> data(esoph)
> myarray <- table(esoph$agegp, esoph$alcgp, esoph$tobgp)
```

- (a) What value is in the myarray cell representing over age 75, consuming 120+ g of alcohol/day and 30+ g of tobacco/day? (Hint: age 75+ is row 6, 120+ g is column 4, and 30+ g is depth slice 4)

```
> myarray
```

```
, , = 0-9g/day
```

	0-39g/day	40-79	80-119	120+
25-34	1	1	1	1
35-44	1	1	1	1
45-54	1	1	1	1
55-64	1	1	1	1
65-74	1	1	1	1
75+	1	1	1	1

```
, , = 10-19
```

	0-39g/day	40-79	80-119	120+
25-34	1	1	1	1
35-44	1	1	1	1
45-54	1	1	1	1
55-64	1	1	1	1
65-74	1	1	1	1
75+	1	1	1	1

```
, , = 20-29
```

	0-39g/day	40-79	80-119	120+
25-34	1	1	0	1
35-44	1	1	1	1
45-54	1	1	1	1
55-64	1	1	1	1
65-74	1	1	1	1

```

75+          0      1      0      0

, , = 30+

```

	0-39g/day	40-79	80-119	120+
25-34	1	1	1	1
35-44	1	1	1	0
45-54	1	1	1	1
55-64	1	1	1	1
65-74	1	0	1	1
75+	1	1	0	0

```
> # 0
```

- (b) How would you index myarray to return that value?

```
> myarray[6,4,4]
[1] 0
```

9. (a) Create a list with two items: a vector with the numbers from 1 to 5 named 'numbers' and a character string with your name named 'name'. Your list should look like:

```
$numbers
[1] 1 2 3 4 5
```

```
$name
[1] "Steve"
```

(Except that your name probably isn't Steve)

```
> mylist <- list(numbers=1:5, name="Steve")
> mylist
```

- (b) How would you index the number 3?

```
> mylist$numbers[3]
[1] 3
```

10. Create a logical vector with the value TRUE for every row in the esoph data frame (that you loaded in step 1) for which ncases is greater than zero. Your vector should look like:

```

[1] FALSE FALSE FALSE FALSE FALSE FALSE
[8] FALSE FALSE FALSE FALSE FALSE TRUE FALSE
[15] FALSE FALSE TRUE FALSE FALSE FALSE TRUE
[22] TRUE FALSE FALSE FALSE FALSE FALSE TRUE
[29] FALSE TRUE TRUE FALSE FALSE FALSE TRUE
[36] TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[43] TRUE TRUE TRUE TRUE TRUE TRUE TRUE

```

```

[50] TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[57] TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[64] TRUE TRUE FALSE TRUE TRUE TRUE TRUE
[71] TRUE TRUE TRUE TRUE TRUE TRUE TRUE
[78] TRUE TRUE TRUE TRUE TRUE FALSE TRUE
[85] TRUE TRUE TRUE TRUE

```

Hint: You can use the `names()` function to return the names of the variables in a data frame

```

> data(esoph)
> mylogical <- esoph$ncases > 0
> mylogical

```

11. create a numerical vector with the number of controls for every row in `esoph` for which there are any cases. Your vector should look like:

```

[1] 1 14 23 14 3 4 46 38 21 15 7 16 14 5 4
[16] 4 4 3 4 49 22 12 6 40 21 17 6 18 15 6
[31] 4 10 7 3 6 48 14 7 34 10 9 13 12 3 1
[46] 4 2 1 1 18 6 3 5 3 1 1 1 2 1

```

```

> control_count <- esoph$ncontrols[esoph$ncases > 0]
> control_count

```

12. Create a 2x2 matrix named `results` with 12 in the a cell, 24 in the b cell, 20 in the c cell and 8 in the d cell as follows:

```

> results <- matrix(c(12, 24, 20, 8), nrow=2, ncol=2, byrow=TRUE)

```

- (a) Calculate the row and column sums using `rowSums()` and `colSums()`

```

> rowSums(results)

```

```

[1] 36 28

```

```

> colSums(results)

```

```

[1] 32 32

```

- (b) Calculate the row and column sums using `apply()`

```

> # Row

```

```

> apply(results, 1, sum)

```

```

[1] 36 28

```

```

> # Column

```

```

> apply(results, 2, sum)

```

```

[1] 32 32

```

- (c) Make a 2-item list named `totals` with first element row sums and second element col sums

```

> totals <- list(rowSums=rowSums(results), colSums=colSums(results))

```

(d) Make a copy of totals named totals.copy

```

> totals.copy <- totals

```

(e) Use mapply to add the row sums and column sums from totals to the row sums and column sums from totals.copy. Your result should look like:

```

rowSums colSums
    128    128

> mapply(sum, totals, totals.copy)

rowSums colSums
    128    128

```

13. First, load the USArrests database built into R using data(USArrests).

```

> data(USArrests)

```

(a) What is the median number of assault arrests per 100,000 people?

```

> median(USArrests$Assault)

[1] 159

```

(b) What was the murder rate in the state(s) that has (have) the median rate of assault arrests have?

```

> USArrests$Murder[USArrests$Assault == median(USArrests$Assault)]

[1] 7.4 4.9

```

(c) which states are they?

```

> rownames(USArrests[which(USArrests$Assault == median(USArrests$Assault)),])

[1] "New Jersey" "Oregon"

```

i. Perform a linear regression predicting the number of murders by the number of assaults.

```

> lm(Murder ~ Assault, data=USArrests)

Call:
lm(formula = Murder ~ Assault, data = USArrests)

Coefficients:
(Intercept)    Assault
    0.63168      0.04191

```

ii. What is the slope of that regression line?

```

> model <- lm(Murder ~ Assault, data=USArrests)
> summary(model)

```

```
Call:
lm(formula = Murder ~ Assault, data = USArrests)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-4.8528 -1.7456 -0.3979  1.3044  7.9256
```

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.631683   0.854776   0.739   0.464
Assault      0.041909   0.004507   9.298 2.6e-12
```

```
(Intercept)
Assault      ***
---

```

```
Signif. codes:
  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 2.629 on 48 degrees of freedom
Multiple R-squared:  0.643,    Adjusted R-squared:  0.6356
F-statistic: 86.45 on 1 and 48 DF,  p-value: 2.596e-12
> # 0.042
```

iii. is it significantly different from 0 at $p < 0.05$?)

```
> summary(model)
```

```
Call:
lm(formula = Murder ~ Assault, data = USArrests)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-4.8528 -1.7456 -0.3979  1.3044  7.9256
```

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.631683   0.854776   0.739   0.464
Assault      0.041909   0.004507   9.298 2.6e-12
```

```
(Intercept)
Assault      ***
---

```

```
Signif. codes:
  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 2.629 on 48 degrees of freedom
Multiple R-squared:  0.643,    Adjusted R-squared:  0.6356
F-statistic: 86.45 on 1 and 48 DF,  p-value: 2.596e-12
```

```
> # Yes
```

14. Duncan questions

- (a) Load the Duncan dataset, first by loading the car package, then loading the dataset itself:

```
> library(car)
> data(Duncan)
> head(Duncan)
```

	type	income	education	prestige
accountant	prof	62	86	82
pilot	prof	72	76	83
architect	prof	75	92	90
author	prof	55	90	76
chemist	prof	64	86	90
minister	prof	21	84	87

- (b) How many of the jobs in the Duncan dataset are type=prof?

```
> table(Duncan$type)
```

bc	prof	wc
21	18	6

```
> #18
```

- (c) Create a new data frame named prof.jobs that is the subset of the Duncan dataset that is professional jobs.

```
> prof.jobs<-subset(Duncan, Duncan$type=='prof')
> prof.jobs
```

	type	income	education	prestige
accountant	prof	62	86	82
pilot	prof	72	76	83
architect	prof	75	92	90
author	prof	55	90	76
chemist	prof	64	86	90
minister	prof	21	84	87
professor	prof	64	93	93
dentist	prof	80	100	90
engineer	prof	72	86	88
undertaker	prof	42	74	57
lawyer	prof	76	98	89
physician	prof	76	97	97
welfare.worker	prof	41	84	59
teacher	prof	48	91	73
contractor	prof	53	45	76
factory.owner	prof	60	56	81
store.manager	prof	42	44	45
banker	prof	78	82	92

- (d) What is the slope of the regression line predicting income from prestige among professional jobs?

```
> x <-lm(income~prestige, data=prof.jobs)
> summary(x)
```

Call:

```
lm(formula = income ~ prestige, data = prof.jobs)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-43.923	-3.146	0.161	8.159	12.849

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.3234	18.1190	0.018	0.98598
prestige	0.7425	0.2220	3.344	0.00412

(Intercept)

prestige **

Signif. codes:

0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 12.91 on 16 degrees of freedom

Multiple R-squared: 0.4114, Adjusted R-squared: 0.3746

F-statistic: 11.18 on 1 and 16 DF, p-value: 0.004117

```
> #Slope = 0.743
```

15. Air quality questions

- (a) Load the airquality dataset that is built into R

```
> data(airquality)
```

- (b) Create a logical variable in the air quality dataset named niceout, which has the value true when the temperature was above 65 and below 80. (You can pick different temperatures if you prefer it to be warmer or colder out)

```
> airquality$niceout <- airquality$Temp >65 & airquality$Temp<80
```

- (c) Use logistic regression to compute odds ratios of being nice out by month.

```
> nice <- glm(niceout ~ as.factor(Month), family=binomial(logit), data=airquality)
```

- (d) What are the odds of a day in June being nice compared to a day in May?

```
> exp(coef(nice))
```



```

      (Intercept) as.factor(Month)6
      0.93750000      1.39487179
as.factor(Month)7 as.factor(Month)8
      0.07356322      0.50793651
as.factor(Month)9
      1.60000000

```

```

> # The odds of 'nice' day in June is 1.39 times the odds of a nice day in May.

```

- (e) Plot the number of nice days in each month (hint: use `tapply()` to sum the number of nice days per month) Use either basic graphics or `ggplot`.

```

> nice_by_month <- data.frame(
+   days=tapply(airquality$niceout, airquality$Month, sum))
> barplot(nice_by_month$days,
+         names.arg=c("May", "June", "July", "August", "September"))

```